



Compute-aware held-out performance (seed 0)

a Accuracy per optimizer update



b Accuracy per recorded training time



Kimi-OPMD uses two optimizer updates per outer iteration; recorded time excludes unlogged overhead.